

Data Manipulation with dplyr

This tutorial is partially based on [Data Analysis and Visualization in R for Ecologists](#) by François Michonneau & Auriel Fournier.

Data for the tutorial is taken from [dslab R package](#)

Introduction

Bracket subsetting is handy, but it can be cumbersome and difficult to read, especially for complicated operations. Enter **dplyr**. **dplyr** is a package for helping with tabular data manipulation. It pairs nicely with **tidyr** which enables you to swiftly convert between different data formats for plotting and analysis.

The **tidyverse** package is an “umbrella-package” that installs **tidyr**, **dplyr**, and several other useful packages for data analysis, such as **ggplot2**, **tibble**, etc.

The **tidyverse** package tries to address 3 common issues that arise when doing data analysis in R:

1. The results from a base R function sometimes depend on the type of data.
2. R expressions are used in a non standard way, which can be confusing for new learners.
3. The existence of hidden arguments having default operations that new learners are not aware of.

To use **tidyverse** package, you can type `install.packages("tidyverse")` straight into the console to install it. Then, type `library(tidyverse)` to load the package.

What is dplyr

dplyr is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges:

- `mutate()` adds new variables that are functions of existing variables
- `select()` picks variables based on their names.

- `filter()` picks cases based on their values.
- `summarise()` reduces multiple values down to a single summary.
- `arrange()` changes the ordering of the rows.

Preparing data

Loading libraries:

```
library(dplyr)
library(dslabs)
```

Preparing and preview data

```
data(gapminder)
gapminder <- gapminder |> as_tibble()
gapminder
```

A tibble: 10,545 x 9

	country	year	infan~1	life_~2	ferti~3	popul~4	gdp	conti~5	region
	<fct>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<fct>	<fct>
1	Albania	1960	115.	62.9	6.19	1.64e6 NA		Europe	South~
2	Algeria	1960	148.	47.5	7.65	1.11e7 1.38e10		Africa	North~
3	Angola	1960	208	36.0	7.32	5.27e6 NA		Africa	Middl~
4	Antigua and Ba~	1960	NA	63.0	4.43	5.47e4 NA		Americ~	Carib~
5	Argentina	1960	59.9	65.4	3.11	2.06e7 1.08e11		Americ~	South~
6	Armenia	1960	NA	66.9	4.55	1.87e6 NA		Asia	Weste~
7	Aruba	1960	NA	65.7	4.82	5.42e4 NA		Americ~	Carib~
8	Australia	1960	20.3	70.9	3.45	1.03e7 9.67e10		Oceania	Austr~
9	Austria	1960	37.3	68.8	2.7	7.07e6 5.24e10		Europe	Weste~
10	Azerbaijan	1960	NA	61.3	5.57	3.90e6 NA		Asia	Weste~

... with 10,535 more rows, and abbreviated variable names
1: infant_mortality, 2: life_expectancy, 3: fertility, 4: population,
5: continent

Inspecting data

```
glimpse(gapminder)
```

Rows: 10,545

Columns: 9

```

$ country      <fct> "Albania", "Algeria", "Angola", "Antigua and Barbuda"~
$ year         <int> 1960, 1960, 1960, 1960, 1960, 1960, 1960, 1960, 1960,~
$ infant_mortality <dbl> 115.40, 148.20, 208.00, NA, 59.87, NA, NA, 20.30, 37.~
$ life_expectancy <dbl> 62.87, 47.50, 35.98, 62.97, 65.39, 66.86, 65.66, 70.8~
$ fertility     <dbl> 6.19, 7.65, 7.32, 4.43, 3.11, 4.55, 4.82, 3.45, 2.70,~
$ population    <dbl> 1636054, 11124892, 5270844, 54681, 20619075, 1867396,~
$ gdp           <dbl> NA, 13828152297, NA, NA, 108322326649, NA, NA, 966778~
$ continent     <fct> Europe, Africa, Africa, Americas, Americas, Asia, Ame~
$ region        <fct> Southern Europe, Northern Africa, Middle Africa, Cari~

```

Selecting columns and filtering rows

To select columns of a data frame, use `select()`. The first argument to this function is the data frame (`gapminder`), and the subsequent arguments are the columns to keep.

```
select(gapminder, country, year, population)
```

```

# A tibble: 10,545 x 3
  country      year population
  <fct>        <int>      <dbl>
1 Albania      1960    1636054
2 Algeria      1960    11124892
3 Angola       1960     5270844
4 Antigua and Barbuda 1960      54681
5 Argentina    1960    20619075
6 Armenia      1960    1867396
7 Aruba        1960     54208
8 Australia    1960    10292328
9 Austria      1960     7065525
10 Azerbaijan  1960     3897889
# ... with 10,535 more rows

```

To select all columns *except* certain ones, put a “-” in front of the variable to exclude it.

```
select(gapminder, -c(gdp, continent, region))
```

```

# A tibble: 10,545 x 6
  country      year infant_mortality life_expectancy fertility popula~1
  <fct>        <int>          <dbl>          <dbl>      <dbl>      <dbl>
1 Albania      1960          115.          62.9        6.19    1636054

```

```

2 Algeria          1960          148.          47.5          7.65 11124892
3 Angola           1960          208           36.0          7.32 5270844
4 Antigua and Barbuda 1960           NA           63.0          4.43 54681
5 Argentina         1960          59.9          65.4          3.11 20619075
6 Armenia           1960           NA           66.9          4.55 1867396
7 Aruba             1960           NA           65.7          4.82 54208
8 Australia         1960          20.3          70.9          3.45 10292328
9 Austria           1960          37.3          68.8          2.7 7065525
10 Azerbaijan       1960           NA           61.3          5.57 3897889
# ... with 10,535 more rows, and abbreviated variable name 1: population

```

This will select all the variables in `gapminder` except `gdp`, `continent` and `region`.

To choose rows based on a specific criterion, use `filter()`:

```
filter(gapminder, year == 2016)
```

```
# A tibble: 185 x 9
```

	country <fct>	year <int>	infan~1 <dbl>	life_~2 <dbl>	ferti~3 <dbl>	popul~4 <dbl>	gdp <dbl>	conti~5 <fct>	region <fct>
1	Albania	2016	NA	78.1	NA	NA	NA	Europe	South~
2	Algeria	2016	NA	76.5	NA	NA	NA	Africa	North~
3	Angola	2016	NA	60	NA	NA	NA	Africa	Middl~
4	Antigua and Barbu~	2016	NA	76.5	NA	NA	NA	Americ~	Carib~
5	Argentina	2016	NA	76.7	NA	NA	NA	Americ~	South~
6	Armenia	2016	NA	74.9	NA	NA	NA	Asia	Weste~
7	Aruba	2016	NA	75.8	NA	NA	NA	Americ~	Carib~
8	Australia	2016	NA	82.3	NA	NA	NA	Oceania	Austr~
9	Austria	2016	NA	81.4	NA	NA	NA	Europe	Weste~
10	Azerbaijan	2016	NA	73.3	NA	NA	NA	Asia	Weste~

```

# ... with 175 more rows, and abbreviated variable names 1: infant_mortality,
# 2: life_expectancy, 3: fertility, 4: population, 5: continent

```

Pipes

What if you want to select and filter at the same time? There are three ways to do this: use intermediate steps, nested functions, or pipes.

With intermediate steps, you create a temporary data frame and use that as input to the next function, like this:

```
df_1 <- filter(gapminder, continent == "Europe")
gapminder_sml <- select(df_1, country, year, gdp)
```

This is readable, but can clutter up your workspace with lots of objects that you have to name individually. With multiple steps, that can be hard to keep track of.

You can also nest functions (i.e. one function inside of another), like this:

```
surveys_sml <- select(filter(gapminder, continent == "Europe"), country, year, gdp)
```

This is handy, but can be difficult to read if too many functions are nested, as R evaluates the expression from the inside out (in this case, filtering, then selecting).

The last option, *pipes*. Pipes let you take the output of one function and send it directly to the next, which is useful when you need to do many things to the same dataset. Pipes in R look like `|>` or `%>%` with **dplyr**. If you use RStudio, you can type the pipe with Ctrl + Shift + M if you have a PC or Cmd + Shift + M if you have a Mac.

```
gapminder |>
  filter(continent == "Europe") |>
  select(country, year, gdp)
```

```
# A tibble: 2,223 x 3
```

	country <fct>	year <int>	gdp <dbl>
1	Albania	1960	NA
2	Austria	1960	52392699681
3	Belarus	1960	NA
4	Belgium	1960	68236665814
5	Bosnia and Herzegovina	1960	NA
6	Bulgaria	1960	NA
7	Croatia	1960	NA
8	Czech Republic	1960	NA
9	Denmark	1960	52164745342
10	Estonia	1960	NA

```
# ... with 2,213 more rows
```

If we want to create a new object with this smaller version of the data, we can assign it a new name:

```
gapminder_sml <- gapminder |>
  filter(continent == "Europe") |>
```

```

    select(country, year, gdp)

gapminder_sml

# A tibble: 2,223 x 3
  country      year      gdp
  <fct>      <int>    <dbl>
1 Albania    1960      NA
2 Austria    1960 52392699681
3 Belarus    1960      NA
4 Belgium    1960 68236665814
5 Bosnia and Herzegovina 1960      NA
6 Bulgaria   1960      NA
7 Croatia    1960      NA
8 Czech Republic 1960      NA
9 Denmark    1960 52164745342
10 Estonia   1960      NA
# ... with 2,213 more rows

```

Note that the final data frame is the leftmost part of this expression.

Mutate

Frequently you'll want to create new columns based on the values in existing columns, for example to do unit conversions, or to find the ratio of values in two columns. For this we'll use `mutate()`.

To create a new column of weight in kg:

```

gapminder %>%
  mutate(population_mln = round(population / 1000000, 2)) |>
  select(country, year, population, population_mln)

```

```

# A tibble: 10,545 x 4
  country      year population population_mln
  <fct>      <int>    <dbl>         <dbl>
1 Albania    1960    1636054         1.64
2 Algeria    1960   11124892        11.1
3 Angola     1960    5270844         5.27
4 Antigua and Barbuda 1960     54681         0.05

```

```

5 Argentina      1960  20619075      20.6
6 Armenia        1960   1867396       1.87
7 Aruba          1960    54208        0.05
8 Australia      1960  10292328      10.3
9 Austria        1960   7065525       7.07
10 Azerbaijan    1960   3897889       3.9
# ... with 10,535 more rows

```

Split-apply-combine data analysis with the `group_by()` and `summarize()` functions

Many data analysis tasks can be approached using the *split-apply-combine* paradigm: split the data into groups, apply some analysis to each group, and then combine the results. Key functions of `dplyr` for this workflow are `group_by()` and `summarize()`.

`group_by()` is often used together with `summarize()`, which collapses each group into a single-row summary of that group. `group_by()` takes as arguments the column names that contain the **categorical** variables for which you want to calculate the summary statistics. So to compute the mean population by country:

```

gapminder |>
  group_by(country) |>
  summarize(mean_population = mean(population, na.rm = TRUE)) |>
  mutate_if(is.numeric, round, 0)

```

```

# A tibble: 185 x 2
  country      mean_population
  <fct>          <dbl>
1 Albania      2708629
2 Algeria      24231378
3 Angola       11909433
4 Antigua and Barbuda 71053
5 Argentina    31638376
6 Armenia      2925011
7 Aruba         74148
8 Australia    16601155
9 Austria       7800180
10 Azerbaijan   6897604
# ... with 175 more rows

```

Counting

When working with data, we often want to know the number of observations found for each factor or combination of factors. For this task, **dplyr** provides `count()`. For example, if we wanted to count the number of rows of data for each country, we would do:

```
gapminder %>%  
  count(country)
```

```
# A tibble: 185 x 2  
  country          n  
  <fct>         <int>  
1 Albania         57  
2 Algeria         57  
3 Angola          57  
4 Antigua and Barbuda 57  
5 Argentina       57  
6 Armenia         57  
7 Aruba           57  
8 Australia       57  
9 Austria         57  
10 Azerbaijan     57  
# ... with 175 more rows
```