Tidyverse. Dplyr

Dr.Sc. Oleksii Yehorchenkov

Department of Spatial Planning

Reading data

```r
1  library(dplyr)
2  library(readr)
3
4  msleep <- read_csv("./data/msleep.csv", show_col_types = FALSE)
```

Data. Mammals sleep

Data: The **msleep** (mammals sleep) data set contains the sleeptimes and weights for a set of mammals. It is part of the **ggplot2** package. This data set contains 83 rows and 11 variables.

```
1  glimpse(msleep)
```

```
Rows: 83
Columns: 11
$ name         <chr> "Cheetah", "Owl monkey", "Mountain beaver", "Greater
shor…
$ genus        <chr> "Acinonyx", "Aotus", "Aplodontia", "Blarina", "Bos",
"Bra…
$ vore         <chr> "carni", "omni", "herbi", "omni", "herbi", "herbi",
"carn…
$ order        <chr> "Carnivora", "Primates", "Rodentia", "Soricomorpha",
"Art…
$ conservation <chr> "lc", NA, "nt", "lc", "domesticated", NA, "vu", NA,
"dome…
$ sleep_total  <dbl> 12.1, 17.0, 14.4, 14.9, 4.0, 14.4, 8.7, 7.0, 10.1, 3.0,
5…
```

The columns of the dataset:

column name	Description
name	common name
genus	taxonomic rank
vore	carnivore, omnivore or herbivore?
order	taxonomic rank
conservation	the conservation status of the mammal
sleep_total	total amount of sleep, in hours
sleep_rem	rem sleep, in hours
sleep_cycle	length of sleep cycle, in hours
awake	amount of time spent awake, in hours
brainwt	brain weight in kilograms
bodywt	body weight in kilograms

Important dplyr verbs to remember

dplyr verbs	Description
select()	select columns
filter()	filter rows
arrange()	re-order or arrange rows
mutate()	create new columns
summarise()	summarise values
group_by()	allows for group operations in the "split-apply-combine" concept

dplyr verbs in action

The two most basic functions are `select()` and `filter()` which selects columns and filters rows, respectively.

Selecting columns using select()

Select a set of columns: the *name* and the *sleep_total* columns.

```r
1  sleepData <- select(msleep, name, sleep_total)
2  head(sleepData)
```

```
# A tibble: 6 × 2
  name                        sleep_total
  <chr>                             <dbl>
1 Cheetah                            12.1
2 Owl monkey                         17
3 Mountain beaver                    14.4
4 Greater short-tailed shrew         14.9
5 Cow                                 4
6 Three-toed sloth                   14.4
```

To select all the columns except a specific column, use the `"-"` `(subtraction)` operator (also known as negative indexing)

```
1 head(select(msleep, -c("name", "sleep_total")))
```

```
# A tibble: 6 × 9
  genus      vore  order conservation sleep_rem sleep_cycle awake  brainwt
bodywt
  <chr>      <chr> <chr> <chr>            <dbl>       <dbl> <dbl>    <dbl>
<dbl>
1 Acinonyx carni Carn… lc                  NA          NA  11.9 NA
50
2 Aotus      omni  Prim… <NA>              1.8          NA   7     0.0155
0.48
3 Aplodon… herbi Rode… nt                  2.4          NA   9.6 NA
1.35
4 Blarina  omni  Sori… lc                  2.3       0.133   9.1  0.00029
0.019
5 Bos        herbi Arti… domesticated      0.7       0.667  20     0.423   600
6 Bradypus herbi Pilo… <NA>                2.2       0.767   9.6 NA
3.85
```

To select a range of columns by name, use the `":"` (colon) operator

```
1  head(select(msleep, name:order))
```

```
# A tibble: 6 × 4
  name                       genus      vore  order
  <chr>                      <chr>      <chr> <chr>
1 Cheetah                    Acinonyx   carni Carnivora
2 Owl monkey                 Aotus      omni  Primates
3 Mountain beaver            Aplodontia herbi Rodentia
4 Greater short-tailed shrew Blarina    omni  Soricomorpha
5 Cow                        Bos        herbi Artiodactyla
6 Three-toed sloth           Bradypus   herbi Pilosa
```

To select all columns that start with the character string *"sl"*, use the function `starts_with()`

```
1  head(select(msleep, starts_with("sl")))
```

```
# A tibble: 6 × 3
  sleep_total sleep_rem sleep_cycle
        <dbl>     <dbl>       <dbl>
1        12.1        NA          NA
2        17         1.8          NA
3        14.4       2.4          NA
4        14.9       2.3       0.133
5         4         0.7       0.667
6        14.4       2.2       0.767
```

Some additional options to select columns based on a specific criteria include

- `ends_with()` = Select columns that end with a character string
- `contains()` = Select columns that contain a character string
- `matches()` = Select columns that match a regular expression
- `one_of()` = Select columns names that are from a group of names

Selecting rows using `filter()`

Filter the rows for mammals that sleep a total of more than 16 hours.

```
1  filter(msleep, sleep_total >= 16)
```

```
# A tibble: 8 × 11
  name    genus  vore   order conservation sleep_total sleep_rem sleep_cycle
awake
  <chr>   <chr>  <chr>  <chr> <chr>              <dbl>     <dbl>       <dbl>
<dbl>
1 Owl mo… Aotus  omni   Prim… <NA>                  17       1.8          NA     7
2 Long-n… Dasy…  carni  Cing… lc                  17.4       3.1       0.383
6.6
3 North … Dide…  omni   Dide… lc                    18       4.9       0.333    6
4 Big br… Epte…  inse…  Chir… lc                  19.7       3.9       0.117
4.3
5 Thick-… Lutr…  carni  Dide… lc                  19.4       6.6          NA
4.6
6 Little… Myot…  inse…  Chir… <NA>                19.9       2         0.2
4.1
7 Giant   Pria…  inse…  Cing… en                  18.1       6.1          NA
```

Filter the rows for mammals that sleep a total of more than 16 hours and have a body weight of greater than 1 kilogram.

```
1  filter(msleep, sleep_total >= 16, bodywt >= 1)
```

```
# A tibble: 3 × 11
  name      genus vore  order conservation sleep_total sleep_rem sleep_cycle
awake
  <chr>     <chr> <chr> <chr> <chr>               <dbl>     <dbl>       <dbl>
<dbl>
1 Long-n... Dasy... carni Cing... lc                 17.4       3.1       0.383
6.6
2 North ... Dide... omni  Dide... lc                 18         4.9       0.333    6
3 Giant ... Prio... inse... Cing... en               18.1       6.1          NA
5.9
# i 2 more variables: brainwt <dbl>, bodywt <dbl>
```

Filter the rows for mammals in the Perissodactyla and Primates taxonomic order

```
1 filter(msleep, order %in% c("Perissodactyla", "Primates"))
```

```
# A tibble: 15 × 11
   name    genus vore  order conservation sleep_total sleep_rem sleep_cycle
awake
   <chr>   <chr> <chr> <chr> <chr>                <dbl>     <dbl>       <dbl>
<dbl>
 1 Owl m…  Aotus omni  Prim… <NA>                    17       1.8          NA     7
 2 Grivet  Cerc… omni  Prim… lc                      10       0.7          NA    14
 3 Horse   Equus herbi Peri… domesticated           2.9       0.6           1
21.1
 4 Donkey  Equus herbi Peri… domesticated           3.1       0.4          NA
20.9
 5 Patas…  Eryt… omni  Prim… lc                    10.9       1.1          NA
13.1
 6 Galago  Gala… omni  Prim… <NA>                   9.8       1.1        0.55
14.2
 7 Human   Homo  omni  Prim… <NA>                     8       1.9         1.5    16
```

You can use the boolean operators (e.g. `>, <, >=, <=, !=, %in%`) to create the logical tests.

Pipe operator: |>

Before we go any futher, let's introduce the pipe operator: |> . This operator allows you to pipe the output from one function to the input of another function. Instead of nesting functions (reading from the inside to the outside), the idea of of piping is to read the functions from left to right.

> **Note**. There is another pipe operator %>% from **magrittr** package with the same functionality. The base pipe operator was introduced in R 4.1.0 in 2021 and it is recommended for use in your code.

Here's an example you have seen:

```
1  head(select(msleep, name, sleep_total))
```

Now in this case, we will pipe the msleep data frame to the function that will select two columns (*name* and *sleep_total*) and then pipe the new data frame to the function head() which will return the head of the new data frame.

```
1  msleep |>
2      select(name, sleep_total) |>
3      head()
```

```
# A tibble: 6 × 2
  name                      sleep_total
  <chr>                           <dbl>
1 Cheetah                          12.1
2 Owl monkey                       17
3 Mountain beaver                  14.4
4 Greater short-tailed shrew       14.9
5 Cow                               4
6 Three-toed sloth                 14.4
```

Back to dplyr verbs in action

Now that you know about the **pipe operator** (|>), we will use it throughout the rest of this tutorial.

You will soon see how useful the **pipe operator** is when we start to combine many functions.

Arrange or re-order rows using arrange()

To arrange (or re-order) rows by a particular column such as the *taxonomic order*, list the name of the column you want to arrange the rows by

```
1  msleep |>
2     select(name, order, sleep_total) |>
3     arrange(order) |>
4     head()
```

```
# A tibble: 6 × 3
  name       order         sleep_total
  <chr>      <chr>               <dbl>
1 Tenrec     Afrosoricida         15.6
2 Cow        Artiodactyla          4
3 Roe deer   Artiodactyla          3
4 Goat       Artiodactyla          5.3
5 Giraffe    Artiodactyla          1.9
6 Sheep      Artiodactyla          3.8
```

Same as above, except here we filter the rows for mammals that sleep for 16 or more hours instead of showing the head of the final data frame

```
1  msleep |>
2      select(name, order, sleep_total) |>
3      arrange(order, sleep_total) |>
4      filter(sleep_total >= 16)
```

```
# A tibble: 8 × 3
  name                  order           sleep_total
  <chr>                 <chr>                 <dbl>
1 Big brown bat         Chiroptera             19.7
2 Little brown bat      Chiroptera             19.9
3 Long-nosed armadillo  Cingulata              17.4
4 Giant armadillo       Cingulata              18.1
5 North American Opossum Didelphimorphia       18
6 Thick-tailed opposum  Didelphimorphia        19.4
7 Owl monkey            Primates               17
8 Arctic ground squirrel Rodentia              16.6
```

Something slightly more complicated: same as above, except arrange the rows in the *sleep_total* column in a descending order. For this, use the function `desc()`

```
1  msleep |>
2      select(name, order, sleep_total) |>
3      arrange(order, desc(sleep_total)) |>
4      filter(sleep_total >= 16)
```

```
# A tibble: 8 × 3
  name                  order             sleep_total
  <chr>                 <chr>                   <dbl>
1 Little brown bat      Chiroptera               19.9
2 Big brown bat         Chiroptera               19.7
3 Giant armadillo       Cingulata                18.1
4 Long-nosed armadillo  Cingulata                17.4
5 Thick-tailed opposum  Didelphimorphia          19.4
6 North American Opossum Didelphimorphia         18
7 Owl monkey            Primates                 17
8 Arctic ground squirrel Rodentia                16.6
```

Create new columns using mutate()

The mutate() function will add new columns to the data frame. Create a new column called *rem_proportion* which is the ratio of rem sleep to total amount of sleep.

```
1  msleep |>
2      mutate(rem_proportion = sleep_rem / sleep_total) |>
3      select(name, sleep_rem, sleep_total, rem_proportion) |>
4      head()
```

```
# A tibble: 6 × 4
  name                     sleep_rem sleep_total rem_proportion
  <chr>                        <dbl>       <dbl>          <dbl>
1 Cheetah                         NA        12.1             NA
2 Owl monkey                     1.8        17            0.106
3 Mountain beaver                2.4        14.4          0.167
4 Greater short-tailed shrew     2.3        14.9          0.154
5 Cow                            0.7         4            0.175
6 Three-toed sloth               2.2        14.4          0.153
```

You can many new columns using `mutate()` (separated by commas). Here we add a second column called *bodywt_grams* which is the *bodywt* column in grams.

```r
1  msleep |>
2      mutate(rem_proportion = sleep_rem / sleep_total,
3             bodywt_grams = bodywt * 1000,
4             wt_ge_50 = if_else(bodywt >= 50, TRUE, FALSE)) |>
5      select(name, sleep_rem, sleep_total,
6             rem_proportion, bodywt, bodywt_grams, wt_ge_50) |>
7      head()
```

```
# A tibble: 6 × 7
  name        sleep_rem sleep_total rem_proportion  bodywt bodywt_grams
wt_ge_50
  <chr>           <dbl>       <dbl>          <dbl>   <dbl>        <dbl> <lgl>
1 Cheetah            NA        12.1             NA      50        50000 TRUE
2 Owl monkey        1.8        17            0.106    0.48          480 FALSE
3 Mountain b…       2.4        14.4          0.167    1.35         1350 FALSE
4 Greater sh…       2.3        14.9          0.154    0.019          19 FALSE
5 Cow               0.7         4            0.175  600          600000 TRUE
6 Three-toed…       2.2        14.4          0.153    3.85         3850 FALSE
```

Create summaries of the data frame using `summarise()`

The `summarise()` function will create summary statistics for a given column in the data frame such as finding the mean. For example, to compute the average number of hours of sleep, apply the `mean()` function to the column *sleep_total* and call the summary value *avg_sleep*.

```
1  msleep |>
2      summarise(avg_sleep = mean(sleep_total))
```

```
# A tibble: 1 × 1
  avg_sleep
      <dbl>
1      10.4
```

There are many other summary statistics you could consider such `sd()`, `min()`, `max()`, `median()`, `sum()`, `n()` (returns the length of vector), `first()` (returns first value in vector), `last()` (returns last value in vector) and `n_distinct()` (number of distinct values in vector).

```
1  msleep |>
2    summarise(avg_sleep = mean(sleep_total),
3             min_sleep = min(sleep_total),
4             max_sleep = max(sleep_total),
5             total = n())
```

```
# A tibble: 1 × 4
  avg_sleep min_sleep max_sleep total
      <dbl>     <dbl>     <dbl> <int>
1      10.4       1.9      19.9    83
```

Group operations using `group_by()`

The `group_by()` verb is an important function in dplyr. As we mentioned before it's related to concept of "split-apply-combine". We literally want to split the data frame by some variable (e.g. taxonomic order), apply a function to the individual data frames and then combine the output.

Let's do that: split the **msleep** data frame by the taxonomic order, then ask for the same summary statistics as above. We expect a set of summary statistics for each taxonomic order.

```
1   msleep |>
2       group_by(order) |>
3       summarise(avg_sleep = mean(sleep_total),
4                 min_sleep = min(sleep_total),
5                 max_sleep = max(sleep_total),
6                 total = n()) |>
7       head()
```

```
# A tibble: 6 × 5
  order         avg_sleep min_sleep max_sleep total
  <chr>             <dbl>     <dbl>     <dbl> <int>
1 Afrosoricida       15.6      15.6      15.6     1
2 Artiodactyla        4.52      1.9       9.1     6
3 Carnivora          10.1       3.5      15.8    12
4 Cetacea             4.5       2.7       5.6     3
5 Chiroptera         19.8      19.7      19.9     2
6 Cingulata          17.8      17.4      18.1     2
```

References

This tutorial is based on

https://genomicsclass.github.io/book/pages/dplyr_tutorial.html